A RESEARCH PROJECT ON

PredictX: Leveraging AI for Intelligent Investment Decisions

Asmita Mondal [2348018]
Sayan Pal [2348056]

PROJECT GUIDE: Dr. Tulasi B

Submitted to the Department of Statistics and Data Science in partial fulfillment of the requirements for the degree of M.Sc. Data Science

TABLE OF CONTENTS

LIVE DEPLOYMENT

https://predictorx.streamlit.app/

Introduction

Background

PredictX is an innovative application designed to simplify and enhance stock analysis. By combining advanced technologies and cutting-edge methodologies, PredictX provides users with an all-in-one platform for stock prediction, visualization, and content generation. Whether you are an individual investor, a financial analyst, or a curious learner, PredictX empowers you with actionable insights and tools to make informed decisions.

How Machine Learning and AI Enhance Stock Analysis

Machine learning (ML) and artificial intelligence (AI) play a pivotal role in transforming stock analysis. By leveraging vast datasets and complex algorithms, these technologies can uncover hidden patterns, provide predictive insights, and enable smarter decision-making. ML and AI are instrumental in analyzing historical trends, forecasting future prices, and generating contextualized content, making them indispensable for modern stock analysis. AI can help leverage these powerful techniques for:

- **Accurate Forecasting:** Predict future stock prices with precision using advanced machine learning models.
- **Data-Driven Insights:** Uncover hidden trends and patterns through sophisticated data analysis.
- **Real-Time Market Monitoring:** Stay updated with real-time market news and data.
- **Personalized Recommendations:** Receive tailored investment advice based on one's risk tolerance and financial goals.

How PredictX Stands Out

Unlike existing stock analysis apps and codes, PredictX integrates the latest technologies and tools into a unified platform to deliver three core functionalities:

1. **Prediction**: Forecast stock prices using advanced machine learning models.
2. **Visualization**: Create dynamic and interactive visual representations of stock data.
3. **Content Generation**: Automatically generate up-to-date financial insights by combining information from The Economic Times and Yahoo Finance YouTube videos.

Empowering Investors

Whether one is a seasoned investor or a novice trader, PredictX provides the tools and knowledge needed to make informed investment decisions. By leveraging the power of AI and machine learning, PredictX empowers one to:

- **Maximize Returns:** Identify profitable investment opportunities.
- **Minimize Risk:** Mitigate potential losses through data-driven insights.
- **Stay Ahead of the Curve:** React quickly to market changes and trends.

With PredictX, one can unlock the full potential of one's investment portfolio and achieve financial goals.

Objectives

- **Visualization**: Allow users to visualize the performance of any selected stock over a customizable date range.
- **Prediction**: Forecast the next day's stock price using five years of historical data, starting from the user-defined date, and provide comparative analysis using five predictive models.
- **Content Generation**: Synthesize the latest stock-related news and insights from The Economic Times and Yahoo Finance YouTube videos, offering users real-time updates.

Data

The data for PredictX is fetched directly using the `yfinance` Python library, which provides a comprehensive and reliable interface for accessing stock market data.

The data for News Content Generation is scraped in real time from The Economic Times news website and Yahoo Finance Youtube Channel @YahooFinance.

Tools and Technologies Used

The following tools and technologies power PredictX, each chosen for its specific utility in stock analysis:

- **Streamlit**: A Python library for creating user-friendly and interactive web applications.
- **Yfinance**: A Python library that allows easy access to historical stock market data.
- **Matplotlib**: A versatile library for creating static, animated, and interactive visualizations.
- **NumPy**: A fundamental library for numerical computing, used for data manipulation and mathematical operations.
- **Pandas**: A data analysis library used for handling and processing stock data efficiently.
- **Seaborn**: A visualization library based on Matplotlib, providing a high-level interface for drawing attractive statistical graphics.
- **Plotly**: An interactive graphing library that enables the creation of dynamic stock visualizations.
- **Python-Dotenv**: Used for managing environment variables securely, such as API keys and configuration settings.
- **Langchain Groq**: A framework that integrates language models with external tools for enhanced AI functionality.
- **Crewal**: A platform for integrating machine learning tools and generating content insights.
- **Crewal Tools**: A suite of tools designed for efficient AI-powered data analysis.
- **Scikit-learn**: A machine learning library for building predictive models.
- **Keras**: A high-level API for building and training neural networks, used in predictive modeling.
- **TensorFlow**: An open-source library for machine learning and deep learning, providing robust support for predictive analytics.
- **Torch**: A machine learning framework for building advanced models, particularly useful in deep learning.
- **Langchain-HuggingFace**: Integrates Hugging Face models for natural language processing tasks, essential for content generation.
- **Streamlit-Lottie**: A library for adding Lottie animations to the Streamlit interface, enhancing the user experience.

About Page


Navigate to
About
Visualize
Predict
News
Ticker Guide
About Our Stock Analysis Tool
Welcome to this crazy shop of stock comprehension!
We're here to present before you a tiny effort towards easement of understanding stocks. This project offers a suite of features made to help you understand and analyze stock data better than us .
Stocks, a dance of rise and fall,
Dreams of fortune or the market's call,
In the gamble of time, we risk it all.
What We Have Built:
1. Visualize :
* Dive deep into historical data for any company of your choice.
* Uncover trends, key metrics, and patterns with interactive visualizations.
* Compare performance against industry peers.

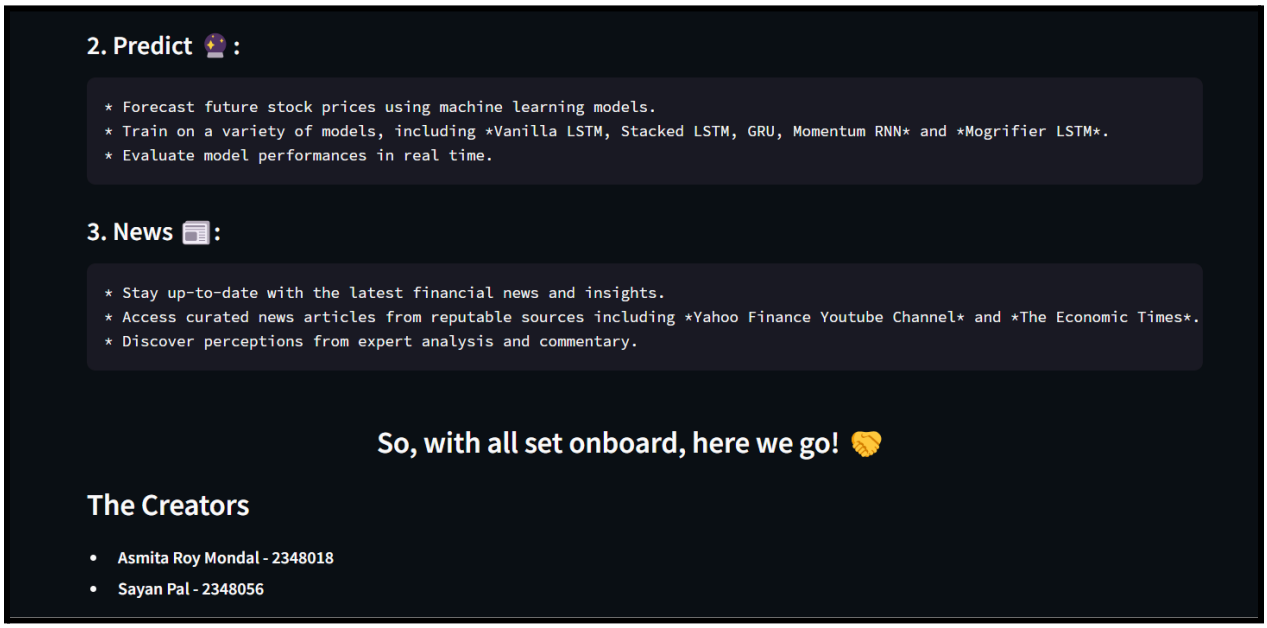
2. Predict :
* Forecast future stock prices using machine learning models.
* Train on a variety of models, including *Vanilla LSTM, Stacked LSTM, GRU, Momentum RNN* and *Mogrifier LSTM*.
* Evaluate model performances in real time.
3. News :
* Stay up-to-date with the latest financial news and insights.
* Access curated news articles from reputable sources including *Yahoo Finance Youtube Channel* and *The Economic Times*.
* Discover perceptions from expert analysis and commentary.
So, with all set onboard, here we go!
The Creators
Asmita Roy Mondal - 2348018
Sayan Pal - 2348056

FEATURE 1: Interactive Real Time Stock Visualization

Overview

The **Stock Comparison Visualization** feature in PredictX allows users to compare the closing prices of multiple stocks over a specified date range. This dynamic and interactive visualization provides a clear and concise representation of stock performance, enabling users to identify trends and patterns across various securities.

Key Functionalities

1. **Stock Selection**:
    - Users can select a primary stock for detailed analysis and add other stocks for comparison.
    - Multiple tickers can be chosen to analyze how different stocks performed during the same timeframe.
2. **Customizable Date Range**:
    - The feature allows users to specify a start and end date, ensuring flexibility in visualizing historical data.
3. **Interactive Graphs**:
    - The visualization is powered by Plotly, offering an interactive experience where users can zoom in/out, hover over data points for specific values, and toggle visibility for selected stocks.
4. **Dynamic Data Integration**:
    - Data is fetched using the `yfinance` library, ensuring that the stock prices displayed are accurate and up-to-date.

Implementation Steps

1. **Fetching Data**:
    - Stock data for the selected tickers is retrieved using the `yfinance` library.
    - The closing prices for each stock within the specified date range are extracted.
2. **Data Aggregation**:
    - The primary stock data is combined with the comparison stocks into a unified dataset.
    - Missing values are handled to ensure consistency in the visualization.
3. **Creating the Plot**:
    - A line chart is created using Plotly, with each stock represented by a distinct line.
    - Interactive features like tooltips, legends, and zoom options enhance user engagement.
4. **Error Handling**:
    - The system gracefully manages cases where:
        - No comparison stocks are selected.
        - Data for the specified tickers or date range is unavailable.

Visuals and Statistics

1. Date Range Selection

- Allows users to specify the time frame for data analysis. The selected range determines the historical stock data retrieved for analysis.

2. Checkbox Options for Metrics

- **Stock Actions:** Displays corporate actions like splits, dividends, etc.
- **Quarterly Financials:** Shows financial statements for the selected stock.
- **Institutional Shareholders:** Highlights major shareholders.
- **Quarterly Cash Flow:** Analyzes cash inflows/outflows.
- **Quarterly Balance Sheet:** Provides an overview of assets, liabilities, and equity.
- **Analysts Recommendation:** Aggregates analysts' buy/hold/sell ratings.

3. Daily Closing Price and Volume

- **Line Charts:** Visualize daily closing prices and trading volume, helping users understand market trends and activity over the selected date range.

4. Bollinger Bands

- Plots price, 20-day moving average, and bands that indicate volatility:
    - **Upper Band:** SMA + (2 * standard deviation).
    - **Lower Band:** SMA - (2 * standard deviation).
- Helps identify overbought or oversold conditions.

5. Key Statistics

- **Average Daily Return:** Indicates typical daily gains/losses.
- **Annualized Volatility:** Measures risk or price fluctuations.
- **Sharpe Ratio:** Assesses risk-adjusted return.
- **Maximum Drawdown:** Captures the largest drop from a peak.

6. Combined Visualization

- **Candlestick Chart:** Provides price movement (open, high, low, close).
- **Volume Bar Chart:** Indicates trading activity.

7. Dividend History

- Displays dividend payments over time or states if no dividends exist.

8. Earnings and Valuation

- **Earnings Per Share (EPS):** Profits per share, a profitability measure.
- **Price-to-Earnings (PE) Ratio:** Valuation metric comparing price to earnings.

9. Risk Metrics

- **Beta:** Stock's volatility relative to the market.
- **Value at Risk (VaR):** Estimates potential loss at a given confidence level.

10. Moving Averages

- Plots short-term and long-term moving averages, showing trends and crossovers (bullish or bearish signals).

11. Technical Indicators

- **SMA (Simple Moving Average):** Average price over a period.
- **EMA (Exponential Moving Average):** Weighted average emphasizing recent data.
- **MACD and Signal Line:** Momentum indicators to identify trend reversals.

12. RSI (Relative Strength Index)

- Measures overbought (>70) or oversold (<30) conditions.
- Visualized with threshold lines at 30 and 70.

13. Closing Price Comparison

- Compares selected stock's closing price with up to three others, providing insights into relative performance.

The provided feature offers a comprehensive and interactive platform for stock visual analysis, making it invaluable for both novice and experienced investors. By integrating various financial metrics, visualizations, and technical indicators, it enables users to explore stock performance across multiple dimensions. Users can analyze daily closing prices, volumes, moving averages, Bollinger Bands, and risk metrics such as beta, Sharpe ratio, and Value at Risk (VaR). The platform also includes advanced features like RSI analysis, candlestick charts, and comparisons with other stocks, allowing users to make informed investment decisions. With its user-friendly Streamlit interface, it fosters engagement through dynamic charts, customizable date ranges, and multiple metrics selection, making it a versatile tool for comprehensive stock analysis.

Design Overview: A Glimpse

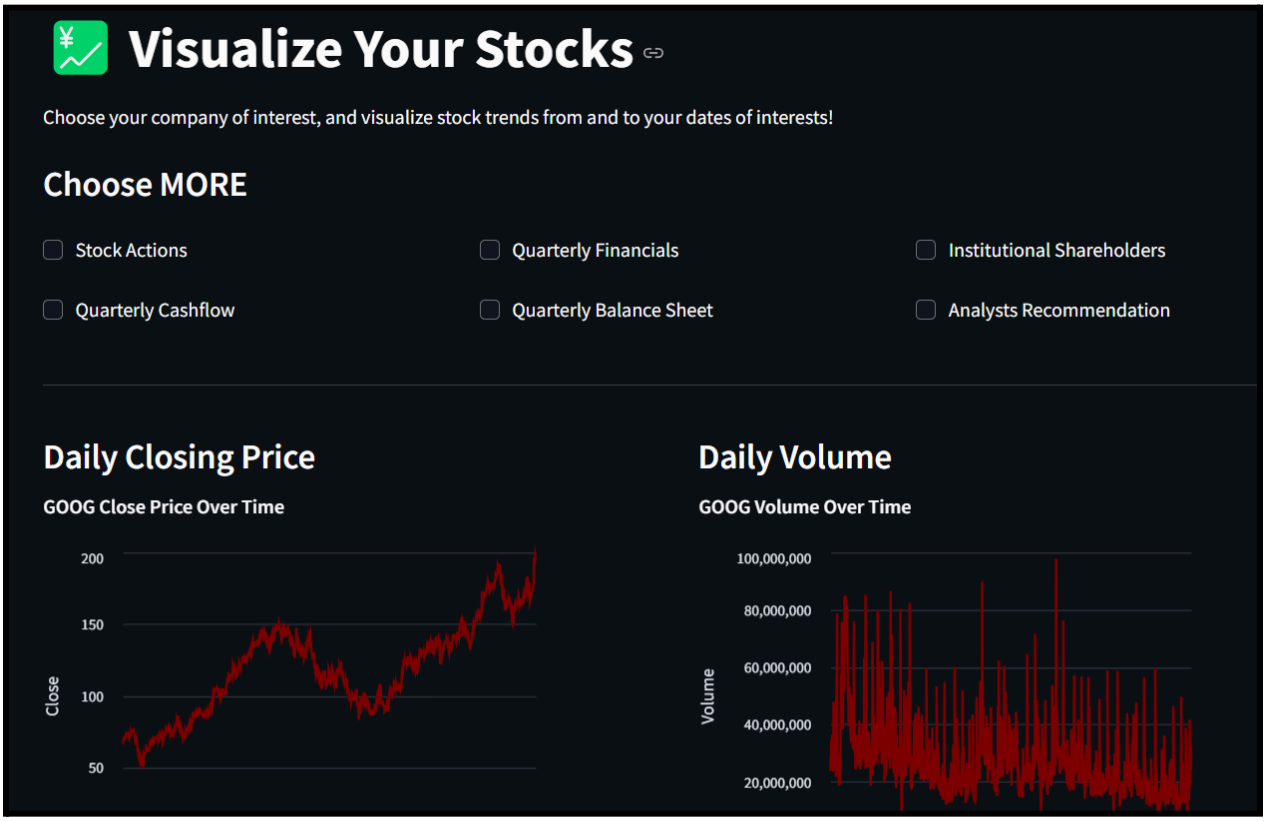

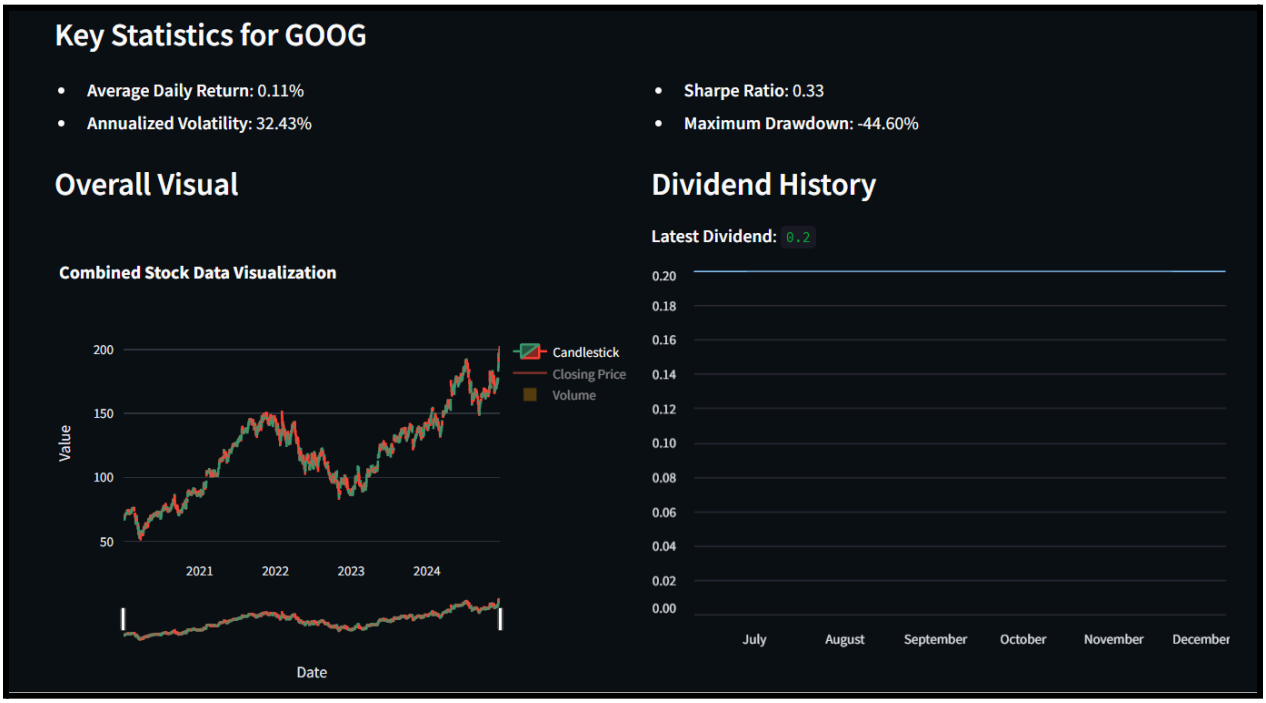

Earnings and Valuation

- **Earnings Per Share (EPS):** 7.55
- **Price-to-Earnings (PE) Ratio:** 26.127151

Risk Metrics

- **Beta:** 1.034
- **95% Value at Risk (VaR):** -3.07%

Moving Averages

Select Moving Average Periods (in days):

20 × 100 ×

Moving Averages for GOOG

RSI Analysis

Select RSI Period:

14

7 28

RSI for GOOG

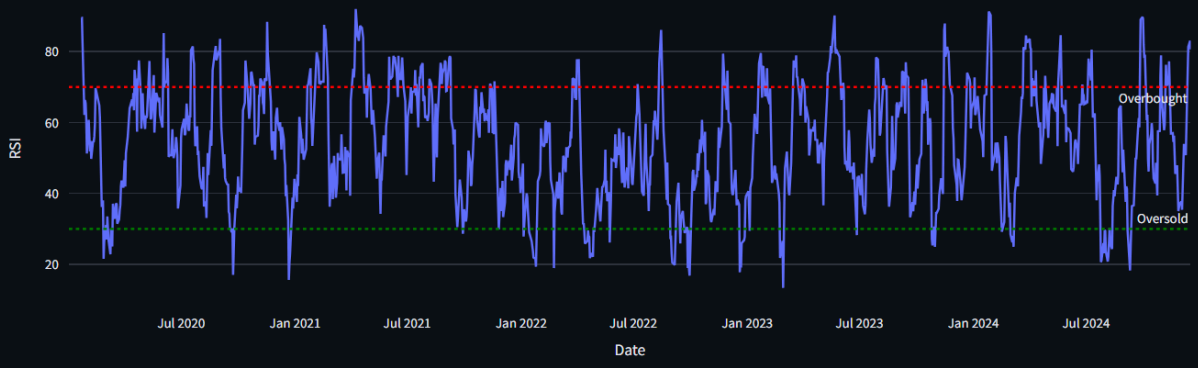

FEATURE 2: Latest Stock Prediction Models' Comparison

Overview

The feature provided is a robust and user-centric implementation for stock market prediction, focusing on historical data visualization and advanced model comparison.

Key Functionalities

1. **Data Fetching and Preprocessing**:
    - The app enables fetching up to five years of historical stock data through `yfinance`, ensuring accuracy with options to manage missing data and outliers.
    - A `MinMaxScaler` is used for scaling data, which is essential for training neural networks effectively.
2. **Interactive Visualization**:
    - Historical stock data is visualized using Plotly candlestick charts.
3. **Model Implementation and User Input**:
    - Five advanced predictive models are built and compared: Vanilla LSTM, Stacked LSTM, GRU, Momentum RNN, and Mogrifier LSTM. Each is tailored for nuanced analysis of stock prices.
    - The architecture of each model includes essential features like dropout layers for regularization, multi-layer configurations, and custom optimizations for improved performance. Epochs are decided by users.
4. **Evaluation and Metrics**:
    - The architecture of each model is displayed. Models are trained and tested with key evaluation metrics (MSE, MAE, RMSE, R^2) displayed in a comparative table for clear insight into their performance.

Models

1. Vanilla LSTM (Long Short-Term Memory)

Architecture:

- **Input Layer**: Takes sequences of stock prices or other time-series data.
- **Single LSTM Layer**: A single LSTM layer with 100 units to capture temporal dependencies in the data.
- **Dropout Layer**: Dropout (0.2) to reduce overfitting by randomly deactivating neurons during training.
- **Output Dense Layer**: A dense layer with a single unit to predict the next stock price.

Why it is Useful:

- LSTM is well-suited for sequential data, such as stock prices, because it can capture long-term dependencies using its memory cell and gates (input, forget, and output gates).
- This architecture is straightforward and computationally less expensive compared to deeper models.
- It is often a good baseline model for time-series prediction tasks.

2. Stacked LSTM

Architecture:

- **Input Layer**: Takes the same sequential input as Vanilla LSTM.
- **Multiple LSTM Layers**:
    - First LSTM layer with 120 units (returning sequences for subsequent layers).
    - Second LSTM layer with 120 units.
    - Third LSTM layer with 100 units.
- **Dropout Layers**: Dropout (0.2 and 0.3) after each LSTM layer to mitigate overfitting.
- **Output Dense Layer**: A single-unit dense layer for prediction.

Why it is Useful:

- Stacked LSTMs are more powerful than a single LSTM layer because they allow the model to learn hierarchical representations of the data.
- The deeper architecture is better for capturing complex patterns and long-term dependencies in time-series data.
- It can lead to higher accuracy but requires more computational power and data to train effectively.

3. GRU (Gated Recurrent Unit)

Architecture:

- **Input Layer**: Similar to LSTM.
- **Single GRU Layer**: A single GRU layer with 100 units.
- **Dropout Layer**: Dropout (0.2) to reduce overfitting.
- **Output Dense Layer**: A single-unit dense layer for prediction.

Why it is Useful:

- GRU is a simpler alternative to LSTM, using fewer gates (update and reset gates).
- It performs similarly to LSTM but is computationally more efficient.
- It is particularly useful when computational resources are limited or the dataset size is small.

4. Momentum RNN

Architecture:

- **Input Layer**: Sequential data input.
- **Simple RNN Layer**: A single Simple RNN layer with 100 units to model temporal dependencies.
- **Dropout Layer**: Dropout (0.2) for regularization.
- **Output Dense Layer**: A single-unit dense layer for prediction.
- **Optimizer**: Uses Momentum-based SGD, which accelerates convergence by adding a fraction of the previous update to the current update.

Why it is Useful:

- Momentum RNN is simpler than LSTM/GRU and focuses on accelerating gradient descent.
- While less powerful than LSTM/GRU, it is efficient for simple sequential tasks.
- Adding momentum helps overcome issues like vanishing gradients in traditional RNNs.

5. Mogrifier LSTM

Architecture:

- **Custom Mogrifier Layer**: A modified LSTM layer that iteratively transforms the input (x) and the hidden state (h) to better capture nonlinear relationships.
    - Includes trainable transformation matrices (`Wx` and `Wh`) applied during multiple rounds of interaction between x and h.
- **Subsequent LSTM Layers**:
    - A standard LSTM layer with 64 units.
- **Dense Layers**:
    - A fully connected layer with 25 units.
    - A final dense layer with 1 unit for prediction.

Why it is Useful:

- Mogrifier LSTM improves the interaction between input and hidden states, leading to better feature extraction and representation.
- It is particularly useful for datasets with complex nonlinear relationships.
- This architecture is more experimental but has the potential to outperform traditional LSTM models for specific tasks.

Metrics

1. Metric 1: R-squared (R^2)

- **Definition**: R^2 measures the proportion of variance in the target variable explained by the model. It ranges from 0 to 1.
- **Interpretation**: A higher R^2 indicates that the model explains a large portion of the variance, with a value close to 1 being ideal. If R^2 is low, the model's predictions are less accurate.

2. Metric 2: Mean Absolute Error (MAE)

- **Definition**: MAE measures the average absolute difference between predicted and actual values. It gives the error in the same units as the target variable.
- **Interpretation**: Lower MAE indicates better model performance, with smaller errors. A high MAE suggests that predictions are not very close to actual values.

3. Metric 3: Mean Squared Error (MSE)

- **Definition**: MSE calculates the average of squared differences between predicted and actual values, penalizing larger errors more than MAE.
- **Interpretation**: A lower MSE indicates better model accuracy, as it reduces the impact of large errors. High MSE signals significant prediction errors.

4. Metric 4: Root Mean Squared Error (RMSE)

- **Definition**: RMSE is the square root of MSE, providing error in the same units as the target variable.
- **Interpretation**: A lower RMSE indicates better model fit, as it penalizes large deviations more heavily. High RMSE indicates the model is less accurate.

Design Overview: A Glimpse

Build Your Stock Model

Choose your company of interest, view a 5 year historical data and watch our models build to give accurate predictions!

Data successfully fetched and preprocessed!

Stock Data for META

	Date	Adj Close META	Close META	High META	Low META	Open META	Volume META
1,251	2024-12-11 00:00:00	632.17	632.68	638.4	621.53	623.37	10,837,200
1,252	2024-12-12 00:00:00	630.2816	630.79	636.66	627.21	631.5	7,474,700
1,253	2024-12-13 00:00:00	619.85	620.35	631.08	616.89	627.22	8,453,300
1,254	2024-12-16 00:00:00	624.24	624.24	631.15	618.58	629.98	10,885,600
1,255	2024-12-17 00:00:00	619.44	619.44	632.37	616.54	626.17	12,887,500

Company Stock Data

META Stock Price Data: 2019-12-20 to 2024-12-17

Training MOMENTUM RNN...

Mogrifier LSTM Model Summary

Model: "sequential_4"

Layer (type)	Output Shape	Param #
custom_mogrifier_lstm (CustomMogrifierLSTM)	(None, 60, 128)	32,768
lstm_5 (LSTM)	(None, 64)	49,408
dense_4 (Dense)	(None, 25)	1,625
dense_5 (Dense)	(None, 1)	26

Total params: 83,827 (327.45 KB)
Trainable params: 83,827 (327.45 KB)
Non-trainable params: 0 (0.00 B)

Training MOGRIFIER LSTM...

All models have been evaluated!

Model Comparisons

	Model	MSE	MAE	RMSE	R2
0	VANILLA LSTM	1,549.3326	35.1	39.3616	0.5465
1	STACKED LSTM	1,861.2116	37.591	43.1418	0.4552
2	GRU	28,339.3093	164.8336	168.3428	-7.2949
3	MOMENTUM RNN	1,405.3531	32.3851	37.488	0.5887
4	MOGRIFIER LSTM	21,468.562	142.4595	146.5215	-5.2839

Best Performing Model: MOMENTUM RNN

R^2 Score: 0.5887

FEATURE 3: Multi Agent Approach to News Content Generation

Overview

The news content generation system scrapes data from reliable sources like The Economic Times news website and Yahoo Finance YouTube channel to gather the latest information based on the user's preferences. The system processes and formats the scraped data into a well-organized, insightful report tailored to the user's interests. By leveraging **CrewAI**, the system utilizes a multi-agent approach, which involves multiple agents working collaboratively to handle different tasks like data extraction, sentiment analysis, and report generation. This multi-agent approach enhances the system's efficiency, scalability, and adaptability, enabling it to process large volumes of information and generate highly relevant, personalized content more effectively than traditional single-agent systems.

Key Functionalities

- **User Input for Topic**: The system prompts users to input a specific stock symbol, company name, or investment topic to generate personalized stock insights.
- **Dynamic Tool Creation**: Based on the user's input, the system creates dynamic tools like **YoutubeChannelSearchTool** (for searching Yahoo Finance YouTube channel) and **ScrapeWebsiteTool** (for scraping news from The Economic Times website) to gather relevant data.
- **Multi-Agent Workflow**: The system uses a multi-agent approach to process data. The agents perform different roles, from research to content generation and formatting, ensuring comprehensive and insightful reports.
- **Insight Generation**: The system generates insights by running sequential tasks across agents, providing a thorough market analysis, transforming the research into actionable content, and finally formatting it into a clean, professional report.
- **Downloadable Report**: After the insights are generated, users can download the report as a `.md` file, which contains the stock insights based on the inputted topic.

Implementation Steps

1. **Tool Setup**:
    - **YoutubeChannelSearchTool** is configured to search Yahoo Finance's YouTube channel for relevant video content.
    - **ScrapeWebsiteTool** is used to scrape news articles and financial reports from The Economic Times website, providing fresh, diverse sources for market analysis.

2. **Dynamic Tool Creation**: The `create_dynamic_tools` function generates the tools based on the user's input, initializing the respective configurations for the Google embedder and Groq LLM models.
3. **Agent Creation and Task Execution**:
    - **Researcher Agent**: The research task focuses on gathering the most recent news and stock market data from multiple sources. It uses Groq's **llama-3.1-8b-instant** model to analyze and process the gathered information.
    - **Insight Compiler Agent**: After research, this agent compiles the findings into a clear and engaging narrative, keeping in mind the readability and actionability of the content.
    - **Formatter Agent**: This agent ensures that the compiled content is well-structured and free of any unneeded elements, maintaining a professional look for the final report.
4. **Multi-Agent Crew**: The `Crew` class brings all agents together, with tasks executed sequentially. The crew's responsibility is to manage the process and ensure the smooth flow of data from one agent to the next, resulting in a polished, informative stock report for the user.

Embedder and Provider

The **embedder providers** used in the system include **Google** and **Groq**. The **Google embedder** uses a pre-trained model (like `embedding-001`) to generate high-quality embeddings from the textual data. This allows the system to perform semantic search and better match the user's query with relevant content. The **Groq embedder**, on the other hand, provides an enhanced language model (LLM) configuration that aids in sophisticated natural language processing tasks like sentiment analysis and insights extraction from unstructured data.

Agents, and Tasks Detail

Researcher Agent (Stock Market Intelligence Analyst):

- **Goal**: To conduct an in-depth investigation and gather the latest information regarding the specified stock or topic.
- **Responsibilities**: This agent is responsible for gathering the most recent and relevant data from the sources provided by the dynamic tools (YouTube and website scraping). It analyzes the latest news, market sentiment, performance trends, and key financial indicators. This agent makes use of the **Groq model** to process and analyze textual data from various sources.
- **Backstory**: The agent is positioned as a seasoned financial analyst with the expertise to evaluate the latest developments in the stock market for the specified company or sector.

Insight Compiler Agent (Financial Content Strategist):

- **Goal**: To transform raw research findings into a coherent, structured, and actionable narrative.
- **Responsibilities**: After the research task is completed, the insights agent takes the raw data and weaves it into a story that provides context, analysis, and strategic insights. This narrative is designed to be clear and professional, ensuring that complex financial data is presented in an easily understandable manner. The agent also utilizes the **Groq model** for summarization and context generation.
- **Backstory**: The agent excels at interpreting complex financial data and presenting it in an engaging, informative style. It ensures that the final output is accessible to users without sacrificing the depth of the analysis.

Formatter Agent (Professional Report Designer):

- **Goal**: To present the generated insights in a polished, well-structured, and professional report.
- **Responsibilities**: The final formatting agent ensures that the financial report is visually clean and well-organized. It removes any placeholder text, unnecessary page breaks, and ensures the report adheres to a strict formatting style for clarity. This agent is crucial in presenting the insights in a readable and professional format. It also uses **Groq model** for final refinements and text formatting.
- **Backstory**: Specializing in document presentation, this agent ensures the content is presented in a manner suitable for formal financial reporting. It provides a well-organized output that is suitable for presentation or further analysis.

Design Overview: A Glimpse

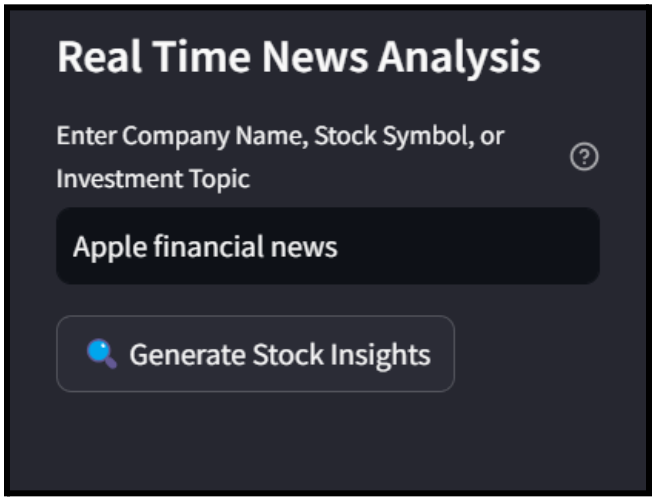

🚀 Generate Stock Insights

Dive deep into personalized stock market insights. Enter a company, stock symbol, or investment topic to get a comprehensive analysis.

✅ Insights Generated Successfully!

📈 Here are the Key Findings

Apple's Record-Breaking First Quarter: A Financial Insights Report

Revenue Growth: The Engine Driving Apple's Success

Apple's revenue growth has been steady over the past five years, with a compound annual growth rate (CAGR) of 10.2%. This growth has been driven by a combination of factors, including:

| Revenue Growth (5-Year CAGR) | 10.2% | | Revenue Breakdown (1Q) | 53% - iPhones, 23% - Macs and iPads, 24% - Services |

Gross and Operating Margins: The Key to Apple's Profitability

Apple's gross margin has increased by 2.5% over the past three years, driven by a combination of factors, including:

| Gross Margin (3-Year Increase) | 2.5% | | Operating Margin (1Q) | 30.2% |

Financial Metrics: A Snapshot of Apple's Strength

Apple's financial metrics paint a picture of a company with a strong balance sheet and a track record of generating significant cash flows. The company's gross margin of 38.2% and operating margin of 29.5% are both impressive, indicating that Apple is able to maintain its pricing power and generate significant profits from its sales.

Metric	Q1 2024	Q1 2023	YoY Change
Revenue	$123.9 billion	$114.5 billion	8.2%
Net Income	$29.1 billion	$23.6 billion	23.4%
Gross Margin	38.2%	37.6%	0.6%
Operating Margin	29.5%	28.4%	1.1%

Market Sentiment: A Buy Recommendation

The market sentiment is overwhelmingly positive, with the average analyst rating for Apple being 2.1 (Buy). This is a testament to the company's strong financial performance and its ability to create long-term value for its shareholders.

Recent Performance Trends: A Story of Volatility

Apple's stock has seen significant volatility in recent months, with a 52-week range of $144.83 to 185.94. While this volatility may be unsettling for some investors, it also presents an opportunity for long-term investors to buy into the company's stock at a discount.

Key Financial Indicators: A Snapshot of Apple's Strength

In conclusion, Apple's Q1 2024 earnings report is a testament to the company's ability to create long-term value for its shareholders. With its strong financial performance, expanding services segment, and commitment to returning capital to shareholders, Apple is poised for continued growth and sustainability in the years to come.

Download Stock Insights Report

Additional Ticker Guide

This lists the 50 most common tickers used in stock market for guiding users in using them for app features.

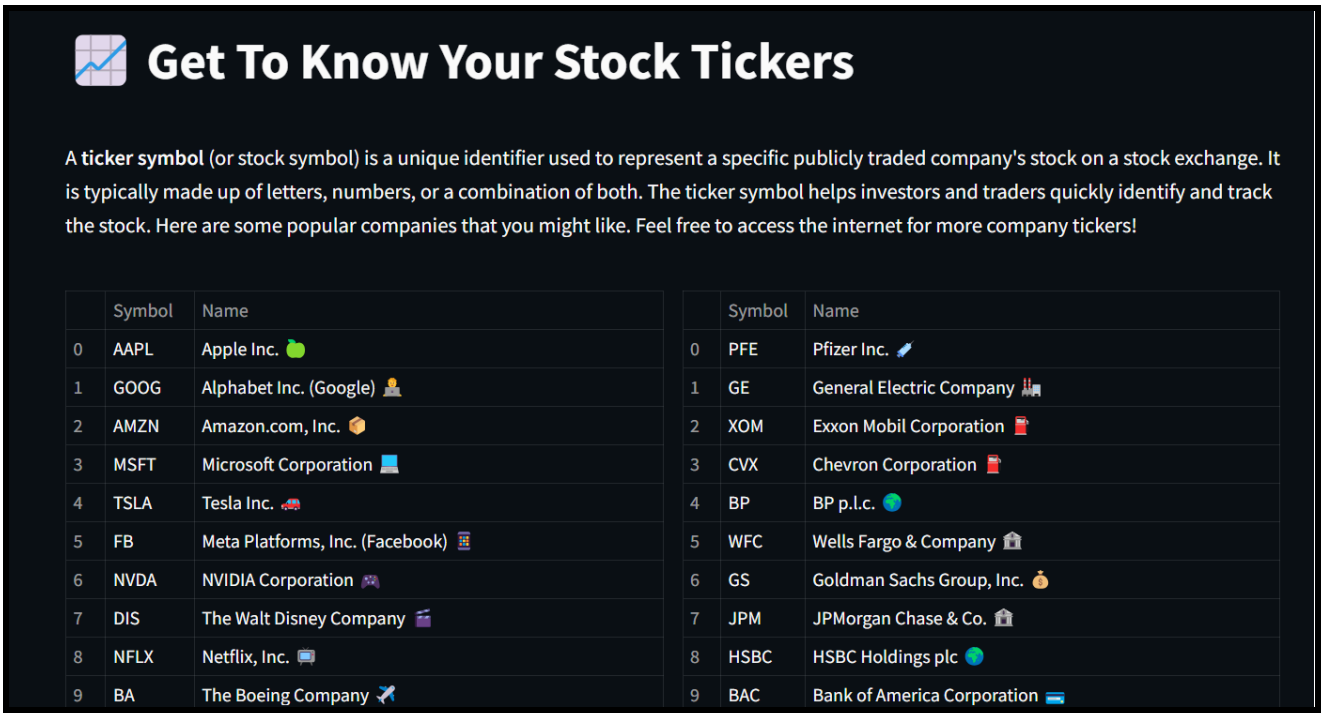

Get To Know Your Stock Tickers

A **ticker symbol** (or stock symbol) is a unique identifier used to represent a specific publicly traded company's stock on a stock exchange. It is typically made up of letters, numbers, or a combination of both. The ticker symbol helps investors and traders quickly identify and track the stock. Here are some popular companies that you might like. Feel free to access the internet for more company tickers!

	Symbol	Name
0	AAPL	Apple Inc.
1	GOOG	Alphabet Inc. (Google)
2	AMZN	Amazon.com, Inc.
3	MSFT	Microsoft Corporation
4	TSLA	Tesla Inc.
5	FB	Meta Platforms, Inc. (Facebook)
6	NVDA	NVIDIA Corporation
7	DIS	The Walt Disney Company
8	NFLX	Netflix, Inc.
9	BA	The Boeing Company

	Symbol	Name
0	PFE	Pfizer Inc.
1	GE	General Electric Company
2	XOM	Exxon Mobil Corporation
3	CVX	Chevron Corporation
4	BP	BP p.l.c.
5	WFC	Wells Fargo & Company
6	GS	Goldman Sachs Group, Inc.
7	JPM	JPMorgan Chase & Co.
8	HSBC	HSBC Holdings plc
9	BAC	Bank of America Corporation

10	WMT	Walmart Inc.
11	JNJ	Johnson & Johnson
12	V	Visa Inc.
13	MA	Mastercard Incorporated
14	CSCO	Cisco Systems, Inc.
15	PYPL	PayPal Holdings, Inc.
16	NVDA	NVIDIA Corporation
17	KO	The Coca-Cola Company
18	PEP	PepsiCo, Inc.
19	INTC	Intel Corporation
20	GM	General Motors Company
21	AMD	Advanced Micro Devices, Inc.
22	ORCL	Oracle Corporation
23	IBM	International Business Machines Corporation
24	GE	General Electric Company
25	T	AT&T Inc.

10	TGT	Target Corporation
11	MCD	McDonald's Corporation
12	SBUX	Starbucks Corporation
13	NKE	Nike, Inc.
14	VZ	Verizon Communications Inc.
15	LYFT	Lyft, Inc.
16	TSM	Taiwan Semiconductor Manufacturing Company
17	SHEL	Shell plc
18	MELI	MercadoLibre, Inc.
19	SHOP	Shopify Inc.
20	BABA	Alibaba Group Holding Limited
21	LULU	Lululemon Athletica Inc.
22	RTX	Raytheon Technologies Corporation
23	SNAP	Snap Inc.
24	TWTR	X (formerly Twitter)
25	SPOT	Spotify Technology S.A.

Theme Change

The entire application can be viewed in **Dark, Light or Custom** themed mode. Here is an example of **LIGHT** theme:

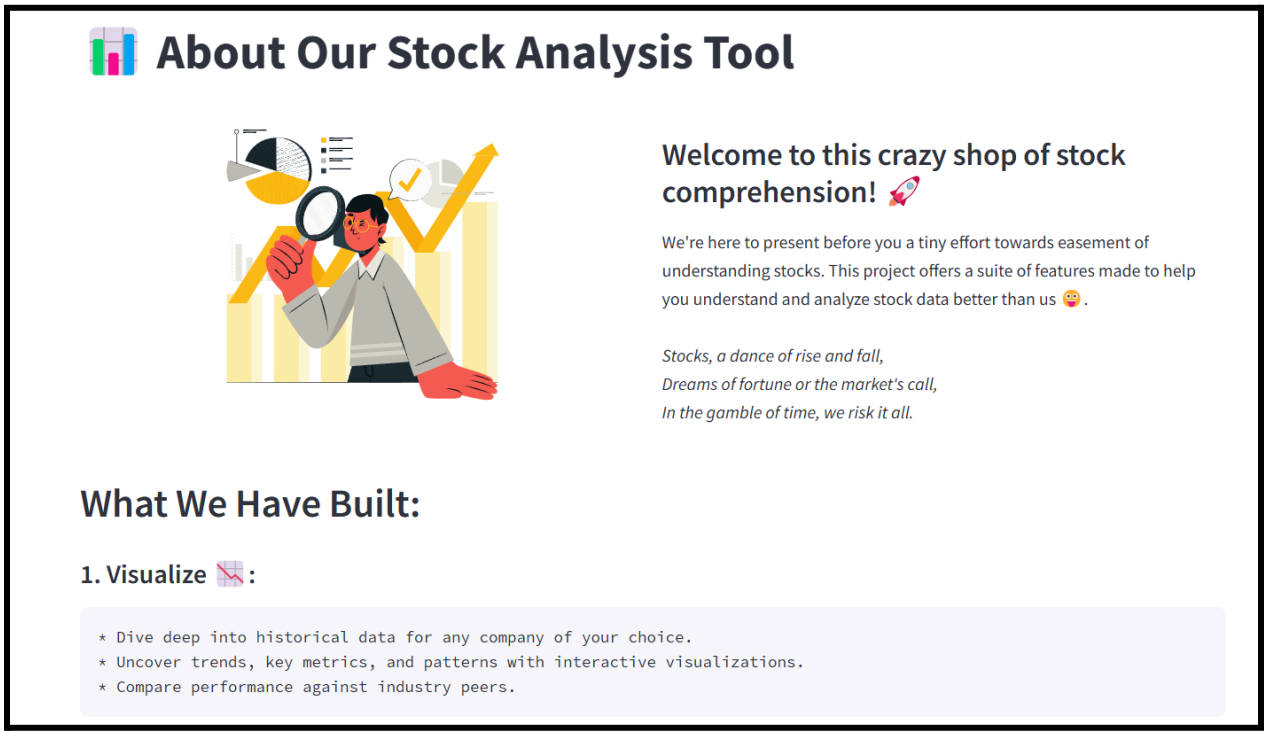

Important Considerations in App Design

Session State Management:

- **Challenge**: Streamlit automatically refreshes the app whenever a widget is clicked, causing the loss of state information and disrupting the flow of interaction.
- **Solution**: To maintain the app's behavior and user input across interactions, **Streamlit's session state** was utilized. This allowed for storing data, such as user inputs and the results of tasks, in memory across widget interactions, ensuring that the app's state is preserved even during refreshes.

Enhanced Visuals with Plotly and Altair:

- **Challenge**: Presenting financial insights in a visually compelling and user-friendly way is essential for engaging users.
- **Solution**: **Plotly** and **Altair** were used to create interactive charts and graphs. Plotly's rich set of interactive visualizations enabled users to explore data through zooming, hovering, and filtering. Altair, known for its declarative syntax, was employed to design clean, concise visual representations of financial data, such as time series graphs and bar charts. These libraries significantly improved the user experience by making complex data more digestible and interactive.

Catchy Lottie Animations:

- **Challenge**: To enhance the visual appeal and keep users engaged, especially when waiting for reports to be generated, a simple yet effective graphic design approach was needed.
- **Solution**: **Lottie animations** were integrated into the app. These lightweight animations added dynamic, eye-catching graphics without heavily impacting load times. Lottie animations helped break the monotony of text-heavy content, drawing users' attention and creating a more engaging and enjoyable experience, especially during the app’s loading or processing stages.

Conclusion

PredictorX app is a powerful tool designed to deliver tailored financial insights by combining cutting-edge AI models and multi-agent systems. It dynamically scrapes data from sources like YouTube and websites, ensuring real-time market insights are provided. The app leverages **Session State** in Streamlit to maintain seamless user interaction, and it incorporates **Plotly** and **Altair** for interactive and enhanced visualizations. To add a layer of engagement, **Lottie animations** were used for smooth and visually appealing graphics.

The app utilizes a range of **prediction models**, including **Vanilla LSTM**, **GRU**, **Bi-directional LSTM**, **GRU with Attention**, and **Transformers**, for analyzing and predicting stock trends. These models were compared based on their ability to handle time series data, capture long-term dependencies, and offer robust forecasting. Additionally, the integration of **Groq models** for natural language processing and **Google Embedders** for advanced semantic analysis ensures that the generated reports are insightful and contextually accurate.

Thus, the app seamlessly combines advanced predictive modeling, web scraping, and AI-driven content generation, providing users with personalized and actionable stock market insights.